

# LAGraph: A Graph Algorithm and Network Analysis Library for GraphBLAS

Timothy A. Davis<sup>\*</sup>, Timothy G. Mattson<sup>†</sup>, James Kitchen<sup>‡</sup>, Scott McMillan<sup>§</sup>, Gábor Szárnyas<sup>¶</sup>, Erik Welch<sup>‡</sup>  
<sup>\*</sup>Texas A&M University <sup>†</sup>Intel Corporation <sup>‡</sup>Anaconda, Inc. <sup>§</sup>Carnegie Mellon University <sup>¶</sup>CWI Amsterdam

**Abstract**—design decisions, easy mode/expert mode, GAP algorithms, ...

**Index Terms**—Graph Processing, Graph Algorithms, Graph Analytics, Linear Algebra, GraphBLAS

## I. INTRODUCTION

LAGraph is a library of Graph Algorithms based on the GraphBLAS

Key contributions:

- document design decisions for LAGraph
- present a concise notation for GraphBLAS algorithms
- algorithms of the GAP benchmark suite [4] used in the IISWC benchmark paper [2]
- improve data ingestion performance, e.g. using SIMD techniques [13]

Recently, numerous graph-specific have targeted GPUs such as Gunrock [18] and GraphBLAST [20], and FPGAs [5].

However, in the near future we expect even more heterogeneous hardware architectures including graph-specific hardware based on the Programmable Integrated Unified Memory Architecture (PIUMA) [1]. Additionally, graph processing workloads can be offloaded to machine learning accelerators, e.g., Tensor Processing Units (TPUs) [11], systolic arrays using reconfigurable dataflow architecture [9], sparse linear algebra-based deep learning accelerators [14].

Previous GraphBLAS design papers: theory [15], C API [17], C++ API [7], distributed API [6], LAGraph [16]

```
int main() {
    return 0; // return zero
}
```

Listing 1: Example

1

## II. DESIGN DECISIONS

Jim

We investigate the following design questions:

- easy/expert mode or standard/advanced more
- multi-threadable
- data structure for representing a graph/matrix
- error handling
- opacity of LAGraph

<sup>1</sup>A non peer-reviewed comparison of 6 popular graph algorithms libraries is available at <https://www.timlrx.com/blog/benchmark-of-popular-graph-network-packages-v2>.




Fig. 1: Separation of concerns using the GraphBLAS API.

- how to work around typecasting being expensive in GraphBLAS
- terminology: properties (parameters? features?)

## III. NOTATION

## IV. ALGORITHMS

GAP algorithms: BFS, SSSP, TC, BC, PR. Not sure whether CC should be included.

### A. BFS

push/pull [19]

GraphBLAST stores both directions and picks automatically [20]

### B. Betweenness Centrality

also does push/pull

### C. PageRank

### D. SSSP

### E. Triangle Count

(Not too interesting from an API design point of view.)

### F. Connected Components

Needs a few GxB extensions.

## V. EVALUATION

Tim D

### A. SuiteSparse Extensions

In a prior paper ([21]), an early draft of SuiteSparse:GraphBLAS v4.0.0 (Aug 2020) was compared with the GAP benchmark [3] and four other graph libraries. This prior version of SS:GrB included two primary data structures for its sparse matrices: compressed sparse vector, and a hypersparse variant [8], both held by row or by column. It included a draft implementation of a bitmap data structure that could only be used in a prototype breadth-first search. Since then, SuiteSparse:GraphBLAS v4.0.3 has been released, with full support for bitmap and full matrices, for all its operations. In an  $m$ -by- $n$  bitmap matrix, the values are held in a full array of size  $mn$ , and another `int8_t` array of size  $mn$  holds the sparsity pattern of the matrix. A full matrix is a simple dense array of size  $mn$ .

The bitmap format is particularly important for the “pull” phase of an algorithm, as used in direction-optimizing breadth-first-search [3]. The GAP benchmark suite uses this method, by holding its frontier as a bitmap in the pull step and a as list in the push step. The GAP BFS was shown to be the fastest BFS amongst the methods compared in [3]. With the addition of the bitmap format to SS:GrB, LAGraph+SS:GrB is able to come within a factor of 2 or so of the performance of the highly-tuned BFS GAP benchmark (see the results in the next section). At the same time, however, the BFS is very easily expressed in an easy-to-read and easy-to-write code, enabling non-experts to obtain a reasonably high level of performance with modest programming effort.

Direction-optimization is incredibly simple to add to an LAGraph algorithm. For example, a batched direction-optimizing betweenness-centrality (BC) algorithm in LAGraph only requires code to determine which direction to use, followed by masked matrix-matrix multiplication  $F(\hat{P}) = FB'$  (the pull) or  $F(\hat{P}) = FA$  (the push), where  $A$  is the adjacency matrix of the graph and  $B = A'$  is its explicit transpose,  $F$  is the frontier, and the complemented mask  $\hat{P}$  is the set of unvisited nodes. The multiplication  $FB'$  relies on the descriptor to represent the transpose of  $B$ , which is not explicitly transposed. In the backward phase, the pull step is  $W = WA'$  while the push is  $W = WB$ , where  $W$  is the 4-by- $n$  matrix in which centrality is accumulated.

Additional optimizations in SS:GrB in the past year include a *lazy sort*. Normally, SS:GrB keeps its vectors sorted (row vectors in a CSR matrix, or column vectors if the matrix is held by column), with entries sorted in ascending order of column or row index, respectively. This simplifies the many algorithms that operate on a `GrB_Matrix`. However, some algorithms naturally produce a jumbled result (matrix multiply in particular), while many algorithms are tolerant of jumbled input matrices. We thus allow the sort to be left pending. The lazy sort joins two other kinds of pending work in SS:GrB: *pending tuples*, and *zombies*. A pending tuple is an entry that is held inside a matrix in an unsorted list, awaiting insertion into the CSR/CSC format of a `GrB_Matrix`. A zombie is the opposite: it is an entry in the CSR/CSC format that has

Algorithm	graph, with run time in seconds				
	Kron	Urand	Twitter	Web	Road
BC GAP	31.52	46.36	10.82	3.01	1.50
BC SS	26.85	31.78	10.04	9.25	51.91
BFS GAP	.31	.58	.22	.34	.25
BFS SS	.52	1.31	.33	.67	3.33
PR GAP	19.81	25.29	15.16	5.13	1.01
PR SS	21.96	27.75	17.22	9.30	1.34
CC GAP	.53	1.66	.23	.22	.05
CC SS	3.42	4.59	1.48	1.97	1.00
SSSP GAP	4.91	7.23	2.02	.81	.21
SSSP SS	17.62	25.62	8.44	9.67	48.49
TC GAP	374.08	21.83	79.58	22.18	.03
TC SS	943.47	34.10	242.36	35.15	.29

TABLE I: Run time of GAP and LAGraph+SS:GrB

been marked for deletion, but has not yet been deleted from the matrix. With the lazy sort, the sort is postponed until another algorithm requires sorted input matrices. If the sort is lazy enough, it might never occur, which is the case for the LAGraph BFS.

Another useful addition to SS:GrB is the new positional binary operators. The BFS relies on the ANY-SECONDI semiring to compute a single step,  $q(\hat{\pi}) = q'A$ , where  $q$  is the current frontier  $\pi$  is the parent vector, and  $A$  is the adjacency matrix.

Consider a matrix multiply for conventional linear algebra, where the PLUS monoid sums a set of  $t$  entries to obtain a single scalar for computing  $c_{ij} = \sum a_{ik}b_{kj}$  in the matrix multiply  $C = AB$ . The ANY monoid performs the reduction of  $t$  entries to a single number by merely selecting any one of the  $t$  entries as the result  $c_{ij}$ . The selection is done non-deterministically, allowing for a benign race condition. In the BFS, this corresponds to selecting any valid parent of a newly discovered node. Indeed, the creation of the ANY operator was inspired by Scott Beamer’s `bfs.cc` method in the GAP benchmark, which has the same benign race condition.

The SECONDI operator is the multiplicative operator in the ANY-SECONDI semiring, where the result of  $a_{ik}b_{kj}$  is simply the index  $k$  in the semiring for  $C = AB$ . This gives the id of the parent node for a newly discovered node in the next frontier.

### B. Performance Results

Our benchmark environment is an NVIDIA DGX Station (donated to Texas A&M by NVIDIA for this research). It includes a 20-core Intel(R) Xeon(R) CPU E5-2698 v4 @ 2.20GHz, with 40 threads, with gcc 5.4.0 -O3. All default settings were used, which means that hyperthreading was enabled. It is a single socket system (no NUMA effects) with 256GB RAM. LAGraph (Jan 29, 2021) and SuiteSparse:GraphBLAS 4.0.3 (Jan 19, 2021) were used. The NVIDIA DGX Station includes four P100 GPUs, but no GPUs were used by this experiment (a GPU-accelerated SS:GrB is in progress). Table I lists the run time (in seconds) for the GAP benchmark and LAGraph+SS:GrB for the 6 algorithms on the 5 benchmark matrices.

With the simple addition of the bitmap (needed for the pull step), the push/pull optimization in BC resulted in a nearly 2x performance gain in the GraphBLAS method for the largest matrices, as compared to the SS:GrB version used for the results presented in [21].

With this change the BC method in LAGraph+SS:GrB is not only expressible in a simple, elegant code, but it is also faster than the highly-tuned GAP benchmark method, `bc.cc`, for the three largest matrices. As compared with the other 4 libraries in [3], this means that LAGraph+SS:GrB is significantly faster than Galois for the Kron, Urand, and Twitter graphs, and on par with the BC methods of GKC and NWGraph. The fourth package, GraphIt, had the fastest BC times for Kron and Urand, at roughly 2.2x and 2.7x faster than the GAP for Kron and Urand, respectively. Of the five packages, GraphIt had the fastest BC time on Twitter, being just barely faster than GAP (1.06x faster). The results in Table I above show that LAGraph+SS:GrB is 1.08x faster than the GAP, on a different machine, putting it on par with GraphIt for this matrix.

We expect the BC in LAGraph+SS:GrB to become faster still in the next release. It does not yet fully exploit the lazy sort. The frontier matrix  $F$  is left jumbled by the lazy sort, but it is sorted right away by a subsequent assignment. Most uses of `GrB_assign` are intolerant of jumbled input matrices, so it sorts them on input. However, `GrB_assign` includes about 40 internal variations, some of which do not actually require sorted input matrices. The particular method used in `GrB_assign` in the LAGraph BC method is one of those methods, so this would be simple to exploit. With this change, the sort would be so lazy that it would *never* occur. The frontier would be computed, left jumbled, used in subsequent computations, and then recomputed (and thus discarded), just as we currently do in the LAGraph BFS.

With the addition of bitmap format (which makes push/pull optimization very simple to express, and very fast) and the ANY-SECONDI semiring, the BFS of a directed or undirected graph is easily expressed in GraphBLAS, and has a performance that is only about 2x slower than the GAP benchmark. We expect the remaining 2x performance gap arises from two issues:

- 1) The GAP assumes that the graph has fewer than 4 billion nodes and edges, and thus uses 32-bit integers throughout. GraphBLAS is written for larger problems, and thus relies solely on 64-bit integers. This cannot be easily “fixed” in GraphBLAS, but in some sense it is the GAP algorithm that would need to be updated in the future. In the benchmark graphs, two graphs are chosen with almost exactly 4 billion edges. Graphs of current interest in large data science can easily exceed 4 billion nodes and edges [12].
- 2) In GraphBLAS, the BFS must be expressed as two calls. The first computes  $q\langle\bar{\pi}\rangle = q'A$ , and the second updates the parent vector,  $\pi\langle q\rangle = q$ :

```
GrB_vxm (q, pi, NULL, semiring, q, A, GrB_DESC_RSC) ;
GrB_assign (pi, q, NULL, q, GrB_ALL, n, GrB_DESC_S) ;
```

In Scott Beamer’s `bfs.cc`, these two steps are fused, and

the matrix-vector multiplication can write its result directly into the vector `pi`. This could be implemented in a future GraphBLAS library, since the GraphBLAS API allows for a non-blocking mode where work is queued and done later, thus enabling a fusion of these two steps. SS:GrB exploits the non-blocking mode (for its lazy sort, pending tuples, and zombies) but does not *yet* exploit the fusion of `GrB_vxm` and `GrB_assign`. We intend to exploit this in the future.

Note that LAGraph+SS:GrB is quite slow for many algorithms (all but PR) on the Road graph, however. The primary reason for this is the high diameter of the Road graph (about 6000). This requires 6000 iterations of GraphBLAS. Each call to GraphBLAS does several `malloc` and `free`s, and in some cases the workspace must be initialized. A future version of SS:GrB is planned that will eliminate this work entirely, by exploiting an internal memory pool. There may be other overheads, but we hope that a memory pool, fusion to fully exploit non-blocking mode, and other optimization will address this large performance gap for the Road graph for these algorithms.

LAGraph+SS:GrB is also up to 3x slower than the GAP for the triangle counting problem (for all but the Road graph, where it is still slower). This can performance gap can be eliminated entirely in the future, if the `GrB_mxm` and `GrB_reduce` are combined in a single fused step, by a full exploitation of the GraphBLAS non-blocking mode. The current method computes  $C\langle L\rangle = LU'$ , followed by the reduction of  $C$  to a single scalar. It would be quite simple to create a fused kernel that does not explicit instantiate the temporary matrix  $C$ .

## VI. UTILITY FUCTIONS

`sort, diag/invdia` (?), `equal`

## VII. CONCLUSION

Future work - ideas:

- Create a Python wrapper for LAGraph
- Implement the LDBC Graphalytics benchmark [10]
- Improve data ingestion performance using e.g., SIMD instructions [13]

## ACKNOWLEDGEMENTS

add acks

G. Szárnyas was partially supported by the SQUIREL-GRAPHS NWO project.

## REFERENCES

- [1] S. Ananthakrishnan *et al.*, “PIUMA: Programmable Integrated Unified Memory Architecture,” *CoRR*, vol. abs/2010.06277, 2020. [Online]. Available: <https://arxiv.org/abs/2010.06277>
- [2] A. Azad *et al.*, “Evaluation of graph analytics frameworks using the GAP Benchmark Suite,” in *IEEE*. IEEE, 2020, pp. 216–227. [Online]. Available: <https://doi.org/10.1109/IISWC50251.2020.00029>
- [3] S. Beamer *et al.*, “Direction-optimizing breadth-first search,” in *SC*. IEEE/ACM, 2012. [Online]. Available: <https://doi.org/10.1109/SC.2012.50>
- [4] —, “The GAP Benchmark Suite,” *CoRR*, vol. abs/1508.03619, 2015. [Online]. Available: <http://arxiv.org/abs/1508.03619>
- [5] M. Besta, D. Stanojevic, J. de Fine Licht, T. Ben-Nun, and T. Hoefler, “Graph processing on FPGAs: Taxonomy, survey, challenges,” *CoRR*, vol. abs/1903.06697, 2019. [Online]. Available: <http://arxiv.org/abs/1903.06697>
- [6] B. Brock *et al.*, “Considerations for a distributed GraphBLAS API,” in *GrAPL at IPDPS*. IEEE, 2020, pp. 215–218. [Online]. Available: <https://doi.org/10.1109/IPDPSW50202.2020.00048>
- [7] —, “A roadmap for the GraphBLAS C++ API,” in *GrAPL at IPDPS*. IEEE, 2020, pp. 219–222. [Online]. Available: <https://doi.org/10.1109/IPDPSW50202.2020.00049>
- [8] A. Buluç and J. R. Gilbert, “On the representation and multiplication of hypersparse matrices,” in *IPDPS*. IEEE, 2008, pp. 1–11. [Online]. Available: <https://doi.org/10.1109/IPDPS.2008.4536313>
- [9] G. F. Grohoski, S. J. Luttrell, R. Prabhakar, R. Sivaramakrishnan, and M. K. Shah, “Virtualization of a reconfigurable data processor,” U.S. Patent 20 200 257 643A1, Aug. 13, 2020.
- [10] A. Iosup *et al.*, “LDBC Graphalytics: A benchmark for large-scale graph analysis on parallel and distributed platforms,” *Vldb*, 2016. [Online]. Available: <http://www.vldb.org/pvldb/vol9/p1317-iosup.pdf>
- [11] N. P. Jouppi *et al.*, “In-datacenter performance analysis of a tensor processing unit,” in *ISCA*. ACM, 2017. [Online]. Available: <https://doi.org/10.1145/3079856.3080246>
- [12] J. Kepner, C. Meiners, C. Byun, S. McGuire, T. Davis, W. Arcand, J. Bernays, D. Bestor, W. Bergeron, V. Gadepally, R. Harnasch, M. Hubbell, M. Houle, M. Jones, A. Kirby, A. Klein, L. Milechin, J. Mullen, A. Prout, A. Reuther, A. Rosa, S. Samsi, D. Stetson, A. Tse, C. Yee, and P. Michaleas, “Multi-temporal analysis and scaling relations of 100,000,000,000 network packets,” in *2020 IEEE High Performance Extreme Computing Conference (HPEC)*, 2020, pp. 1–6.
- [13] G. Langdale and D. Lemire, “Parsing gigabytes of JSON per second,” *Vldb J.*, vol. 28, no. 6, pp. 941–960, 2019. [Online]. Available: <https://doi.org/10.1007/s00778-019-00578-5>
- [14] S. Lie, G. R. Lauterbach, M. E. James, M. Morrison, and S. Arekapudi, “Dataflow triggered tasks for accelerated deep learning,” U.S. Patent 10 614 357B2, Apr. 7, 2020.
- [15] T. Mattson *et al.*, “Standards for graph algorithm primitives,” in *HPEC*. IEEE, 2013. [Online]. Available: <https://doi.org/10.1109/HPEC.2013.6670338>
- [16] —, “LAGraph: A community effort to collect graph algorithms built on top of the GraphBLAS,” in *GrAPL at IPDPS*, 2019. [Online]. Available: <https://doi.org/10.1109/IPDPSW.2019.00053>
- [17] T. G. Mattson *et al.*, “GraphBLAS C API: Ideas for future versions of the specification,” in *HPEC*. IEEE, 2017. [Online]. Available: <https://doi.org/10.1109/HPEC.2017.8091095>
- [18] Y. Wang *et al.*, “Gunrock: GPU graph analytics,” *ACM Trans. Parallel Comput.*, vol. 4, no. 1, pp. 3:1–3:49, 2017. [Online]. Available: <https://doi.org/10.1145/3108140>
- [19] C. Yang *et al.*, “Implementing push-pull efficiently in GraphBLAS,” in *ICPP*. ACM, 2018, pp. 89:1–89:11. [Online]. Available: <https://doi.org/10.1145/3225058.3225122>
- [20] —, “GraphBLAST: A high-performance linear algebra-based graph framework on the GPU,” *CoRR*, vol. abs/1908.01407, 2019. [Online]. Available: <http://arxiv.org/abs/1908.01407>
- [21] Y. Zhang *et al.*, “FastSV: A distributed-memory connected component algorithm with fast convergence,” in *PPSC*. SIAM, 2020, pp. 46–57. [Online]. Available: <https://doi.org/10.1137/1.9781611976137.5>